2026 16th International Conference on System Engineering & Technology (ICSET)
3-4 November 2026, Bandung, Indonesia

# Physics-Informed Machine Learning for Parameter Estimation in Rigid Multi-Quadrotor Transport System

Miqdad Abdurrahman, Anggera Bayuwindra, Hilton Tnunay
*Sekolah Teknik Elektro dan Informatika (STEI)*
*Institut Teknologi Bandung (ITB)*
Bandung, Indonesia
Email: 23224011@mahasiswa.itb.ac.id, bayuwindra@itb.ac.id, hilton.tnunay@itb.ac.id

***Abstract*—Cooperative multi-quadrotor systems rigidly attached to a common frame provide scalable solutions for heavy-lift transportation. However, unmodeled payload mass and unknown spatial offsets can introduce continuous disturbance torques, degrading controller performance and causing lateral drift. This paper introduces a fixed time-step simulation of a Physics-Informed Neural Network (PINN) parameter estimation framework to dynamically identify unmodeled, off-center payloads. To achieve online parameter adaptation, the proposed PINN utilizes continuous gradient descent over a short-term receding data horizon. By embedding translational and rotational physics residuals derived from Euler-Lagrange and Euler-Poincaré equations, the network enforces globally valid rigid-body constraints. The estimated parameters are then continuously fed into the system's control allocation matrix, enabling immediate and adaptive thrust reallocation. Simulations confirm that the integrated PINN estimator converges the desired parameters, arrests transient lateral drift, and successfully restores trajectory tracking stability during dynamic flight maneuvers.** [1]

***Index Terms*—Multi-Agent Systems, Physics-Informed Neural Networks, Cooperative Payload Transport, Parameter Estimation, Unmanned Aerial Vehicles.**

## I. Introduction

### A. Motivation and Related Works

Cooperative multi-quadrotor systems provide scalable, high-capacity solutions for payload transportation [1], [2]. While cable-suspended payloads [3], [4] are common, rigidly attaching quadrotors to a common frame allows the distributed system to be controlled as a single composite rigid body [2]. However, this geometric simplification assumes a perfectly symmetrical mass distribution [5], [6]. In practice, uncertain payload Centers of Gravity (CoG) introduce unmodeled disturbance torques, causing severe lateral drift that may lead to a crash [7], [8].

To mitigate these unmodeled disturbances, recent literature has explored structural hardware optimizations, such as canting rotors to artificially increase yaw authority [2], alongside advanced algorithmic adaptations. Various parameter estimation and control architectures have been successfully deployed to handle irregular payloads, including classical batch and recursive least-squares algorithms [5], Concurrent Learning (CL) frameworks [7], and Digital Twin (DT) Model Predictive Control (MPC) architectures [8]. Building upon this established body of work, there is an opportunity to explore emerging machine learning techniques as alternative estimation mechanisms for real-time flight stabilization.

This research develops a Proof of Concept (PoC) for a Physics-Informed Neural Network (PINN) estimator tailored for cooperative, rigid multi-quadrotor payload transportation within a simulated environment. Specifically, this work explores the feasibility of using a PINN to mitigate dynamic disturbances caused by an unmodeled, off-center static payload. To bridge the gap between PINN-based parameter estimation and multi-quadrotor control, the primary contributions of this work are:

1) A PINN estimator that dynamically identifies unknown payload mass and static spatial shifts by utilizing a short-term receding data horizon;
2) Embedding an energy-based loss function derived from Euler-Lagrange and Euler-Poincaré equations;
3) Integrating the PINN estimator with a small-angle multi-agent controller based on [5] to continuously feed updated offset parameters into the control allocation matrix.

## II. Preliminaries

### A. Physics-Informed Parameter Estimation

Physics-Informed Neural Networks (PINNs) represent a paradigm shift in scientific machine learning by seamlessly embedding known physical laws into the deep learning optimization process. Unlike purely data-driven models that rely exclusively on large empirical datasets to learn system behaviors implicitly, PINNs utilize governing Partial Differential Equations (PDEs) or Ordinary Differential Equations (ODEs) as explicit regularizers within their loss function [9], [10].

The network typically consists of fully connected Artificial Neural Network (ANN) layers parameterized by weights and biases $\theta$ [9]. To illustrate the fundamental mechanism, consider

[1]The source code for the PINN estimator and simulation environment is available at: https://github.com/miqdude/pinn_invese_estimator_conf_repo

XXX-X-XXXX-XXXX-X/XX/$XX.00 ©20XX IEEE

a general dynamical system governed by an ODE of the form $\mathcal{F}(u, \dot{u}, \ddot{u}, t) = 0$. The PINN acts as a surrogate model, taking the independent variable time $t$ as its sole input and outputting a predicted system state $\hat{u}(t)$:

$$\hat{u}(t) = \mathcal{N}_\theta(t). \tag{1}$$

During training, the network minimizes a composite loss function $\mathcal{L}_{total}$ that balances empirical observation with physical consistency. The data loss ($\mathcal{L}_{data}$) anchors the network to reality by minimizing the Mean Squared Error (MSE) between the network's predictions $\hat{u}(t)$ measured state $u_{true}$, multiply with a scalar hyperparameter $\lambda_1$ at specific time steps $t_i$:

$$\mathcal{L}_{data} = \lambda_1 \frac{1}{N} \sum_{i=1}^{N} \left\| \hat{u}(t_i) - u_{true}(t_i) \right\|^2. \tag{2}$$

Simultaneously, the network evaluates the physical validity of its predictions. Rather than relying on numerical approximation, PINNs leverage Automatic Differentiation (Autograd) to compute exact analytical derivatives of the output with respect to the input (e.g., $\dot{\hat{u}} = \frac{\partial \mathcal{N}_\theta}{\partial t}$). These derivatives are substituted back into the governing equation to form a physics residual $R(t) = \mathcal{F}(\hat{u}, \dot{\hat{u}}, \ddot{\hat{u}}, t)$. The physics loss ($\mathcal{L}_{physics}$) penalizes any deviation from zero across $M$ collocation points:

$$\mathcal{L}_{physics} = \lambda_2 \frac{1}{M} \sum_{j=1}^{M} \left\| R(t_j) \right\|^2. \tag{3}$$

The total loss is calculated as a weighted sum of these objectives, $\mathcal{L}_{total} = \lambda_1 \mathcal{L}_{data} + \lambda_2 \mathcal{L}_{physics}$. By minimizing this composite function, the network learns to predict states that not only match the data but also strictly obey the underlying physical mechanics of the system [10].

## III. System Description

### A. System Modeling

The rigid-frame vehicle consists of a 0.8 kg, $1 \times 1$ m square board with four 1.5 kg Iris quadrotors modeled by [11], rigidly mounted at its corners, displaced by $\pm 0.5$ m from the geometric center. We define an inertial world frame $\mathcal{W}$ with an upward-pointing Z-axis, and a primary body frame $\mathcal{B}$ anchored at the board's geometric center. Within $\mathcal{B}$, the positive X-axis designates the forward direction of the vehicle, and the positive Y-axis points to the left, completing a right-handed coordinate system. At the identity orientation (level and unrotated), the axes of $\mathcal{B}$ perfectly align with $\mathcal{W}$. Each quadrotor is modeled as an ideal actuator producing direct thrust, and each maintains a local coordinate frame that is strictly parallel to $\mathcal{B}$. Due to the rigid attachments, there is no relative rotation between the drones and the central board; their alignment differs from $\mathcal{B}$ purely by a fixed spatial translation in the body's XY plane. The vehicle's global position in $\mathcal{W}$ is denoted by $\mathbf{r} = [x, y, z]^T$, and its orientation is parameterized by the Euler angles $\eta = [\phi, \theta, \psi]^T$.

The MAS employs a centralized control architecture derived from [5]. The composite rigid body is commanded via a unified control wrench $\boldsymbol{\tau}$, defined as:

$$\boldsymbol{\tau} = [F_B, M_{xB}, M_{yB}, M_{zB}]^T, \tag{4}$$

where $F_B$ represents the total collective thrust acting along the frame's $z$-axis, while $M_{xB}$, $M_{yB}$, and $M_{zB}$ denote the respective desired roll, pitch, and yaw control moments applied about the frame which relate to individual quadrotor thrust inputs $\mathbf{u} = [F_1, F_2, F_3, F_4]^T$ via the allocation matrix $A \in \mathbb{R}^{4\times4}$:

$$\boldsymbol{\tau} = A\mathbf{u}, \tag{5}$$

$$A = \begin{bmatrix} 1 & 1 & 1 & 1 \\ y_1 & y_2 & y_3 & y_4 \\ -x_1 & -x_2 & -x_3 & -x_4 \\ c_1 & c_2 & c_3 & c_4 \end{bmatrix}. \tag{6}$$

The rows of $A$ denote total thrust, roll moment, pitch moment, and yaw reactive torque, respectively.

### B. Dynamic Thrust Allocation

When a payload shifts from the body's central frame, it alters the effective moment arms. Let $\hat{\mathbf{d}}_{\text{cog}} = [\hat{x}_{\text{cog}}, \hat{y}_{\text{cog}}]^T$ denote the estimated offset. Redefining the lever arms as $\Delta x_i = x_i - \hat{x}_{\text{cog}}$ and $\Delta y_i = y_i - \hat{y}_{\text{cog}}$, the allocation matrix adapts dynamically:

$$A_{\text{final}} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ \Delta y_1 & \Delta y_2 & \Delta y_3 & \Delta y_4 \\ -\Delta x_1 & -\Delta x_2 & -\Delta x_3 & -\Delta x_4 \\ c_1 & c_2 & c_3 & c_4 \end{bmatrix}. \tag{7}$$

Control inputs are computed via the Moore-Penrose pseudo-inverse to handle transient rank deficiencies:

$$\mathbf{u} = A_{\text{final}}^{+} \boldsymbol{\tau}. \tag{8}$$

### C. Centralized Control Design

Based on [5], the controller separates position and attitude tracking. Position and velocity errors define commanded acceleration $\ddot{\mathbf{r}}_{cmd}$:

$$\mathbf{e}_p = \mathbf{r}_{des} - \mathbf{r}, \quad \mathbf{e}_v = \dot{\mathbf{r}}_{des} - \dot{\mathbf{r}}, \tag{9}$$

$$\ddot{\mathbf{r}}_{cmd} = K_p \mathbf{e}_p + K_d \mathbf{e}_v + \ddot{\mathbf{r}}_{\text{ff}}. \tag{10}$$

Desired roll ($\phi^{des}$), pitch ($\theta^{des}$), and total thrust ($F_B^{des}$) are extracted via small-angle approximation:

$$\begin{aligned} \phi^{des} &= \frac{1}{g}(\ddot{r}_x^{des} \sin\psi_T - \ddot{r}_y^{des} \cos\psi_T), \\ \theta^{des} &= \frac{1}{g}(\ddot{r}_y^{des} \cos\psi_T + \ddot{r}_y^{des} \sin\psi_T), \\ F_B^{des} &= m(\ddot{r}_z^{des} + g). \end{aligned} \tag{11}$$

Inner-loop desired moments $\mathbf{M}_B^{des}$ regulate orientation against angular velocities $(p, q, r)$:

$$\begin{aligned} M_{xB}^{des} &= k_{p,\phi}(\phi^{des} - \phi) + k_{d,\phi}(p^{des} - p), \\ M_{yB}^{des} &= k_{p,\theta}(\theta^{des} - \theta) + k_{d,\theta}(q^{des} - q), \\ M_{zB}^{des} &= k_{p,\psi}(\psi^{des} - \psi) + k_{d,\psi}(r^{des} - r). \end{aligned} \tag{12}$$

## IV. Proposed PINN-Based Estimator Design

### A. PINN Architecture

To address the 6-DOF rigid-body dynamics, our PINN framework utilizes continuous time $t$ as the sole independent variable to solve the time-dependent ODEs. The output is a 12-dimensional state vector consisting of posisition, linear velocity, attitude, and angular velocity [10], [12], [13]. The PINN design used in this research is described in Table I.

TABLE I
Proposed Fully Connected layer PINN Architecture Specifications

| Layer | Number of Neurons | Activation Function |
|---|---|---|
| Input | 1 (Continuous time $t$) | None |
| Hidden 1 | 64 | tanh |
| Hidden 2 | 64 | tanh |
| Output | 12 (Kinematic state vector) | Linear |

### B. Loss Function Formulation

The neural network acts as a surrogate model mapping continuous time $t$ to the predicted 12 states of the system. We define the complete state vector $\mathbf{s}$, comprising the position, linear velocity, attitude, and angular velocity:

$$\mathbf{s} = [\mathbf{r}^T, \dot{\mathbf{r}}^T, \boldsymbol{\eta}^T, \boldsymbol{\omega}^T]^T. \tag{13}$$

To train the network, we formulate a composite loss function $\mathcal{L}_{total}$ [13], [14] that simultaneously penalizes empirical data deviations and violations of rigid-body dynamics. This objective function is defined as a weighted sum of the data-driven observation loss ($\mathcal{L}_{data}$), the translational physics residual ($\mathcal{L}_{trans}$), and the rotational physics residual ($\mathcal{L}_{rot}$):

$$\mathcal{L}_{total} = \lambda_1 \mathcal{L}_{data} + \lambda_2 \mathcal{L}_{trans} + \lambda_3 \mathcal{L}_{rot}, \tag{14}$$

where the hyperparameters $\lambda_1, \lambda_2$, and $\lambda_3$ are constants to balance the contributions of each loss term.

The data loss ensures the network's predictions ($\hat{s}$) anchor to the physical reality measured by the onboard sensors. It is computed as the Mean Squared Error (MSE) over a temporal window of $N$ collocation points:

$$\mathcal{L}_{data} = \frac{1}{N} \sum_{i=1}^{N} \left\| \hat{s}_i - s_{true,i} \right\|^2. \tag{15}$$

To enforce the translational physical constraints, we formulate a residual loss based on the Euler-Lagrange equations [15]. Let $L = T_{trans} - V$ denote the system Lagrangian, where $T_{trans}$ is the kinetic energy and $V$ is the potential energy. The translational dynamics must satisfy:

$$\frac{d}{dt}\left(\frac{\partial L}{\partial \mathbf{v}}\right) - \frac{\partial L}{\partial \mathbf{r}} = \mathbf{F}_{ext}. \tag{16}$$

Evaluating the kinetic inertial term yields $m\dot{\mathbf{v}}$. Rather than analytically deriving the potential gradient, we leverage PyTorch's Automatic Differentiation (Autograd) engine to dynamically compute $\frac{\partial V}{\partial \mathbf{r}}$ during the forward pass. The resulting translational physics loss minimizes the residual of these forces across the temporal window:

$$\mathcal{L}_{trans} = \frac{1}{N} \sum_{i=1}^{N} \left\| m\dot{\mathbf{v}}_i + \frac{\partial V}{\partial \mathbf{r}_i} - \mathbf{F}_{ext,i} \right\|^2. \tag{17}$$

While translational dynamics are straightforward, evaluating rotational dynamics using minimal coordinates (e.g., Euler angles) introduces mathematical singularities, commonly known as gimbal lock, at extreme pitch or roll angles. To ensure global mathematical validity, we evaluate the rotational dynamics directly on the Special Orthogonal group, $SO(3)$. The rotational kinetic energy of the rigid multi-agent system is defined as:

$$T_{rot} = \frac{1}{2} \boldsymbol{\omega}^T \mathbf{J} \boldsymbol{\omega}, \tag{18}$$

where $\mathbf{J} \in \mathbb{R}^{3\times3}$ is the estimated inertia tensor. By applying the Euler-Poincaré equations of motion for rigid body we get:

$$\frac{d}{dt}\left(\frac{\partial T_{rot}}{\partial \boldsymbol{\omega}}\right) + \boldsymbol{\omega} \times \frac{\partial T_{rot}}{\partial \boldsymbol{\omega}} = \boldsymbol{\tau}_{ext}. \tag{19}$$

The partial derivative $\frac{\partial T_{rot}}{\partial \boldsymbol{\omega}} = \mathbf{J}\boldsymbol{\omega}$ represents the system's angular momentum, Eq. (19) resolves into the Newton-Euler formulation for rigid body rotation. In our cooperative framework, the total external torque $\boldsymbol{\tau}_{ext}$ comprises the known applied control wrench $\boldsymbol{\tau}_{app}$ and the unknown disturbance torque $\boldsymbol{\tau}_{dist}$ generated by the payload's CoG offset. The rotational physics residual evaluated during the forward pass is thus formulated as:

$$\mathcal{R}_{rot} = \left(\hat{\mathbf{J}}\dot{\hat{\boldsymbol{\omega}}} + \hat{\boldsymbol{\omega}} \times (\hat{\mathbf{J}}\hat{\boldsymbol{\omega}})\right) - (\boldsymbol{\tau}_{app} - \boldsymbol{\tau}_{dist}). \tag{20}$$

Finally, the rotational loss penalizes the mean squared error of this residual, forcing the neural network to implicitly map the physical discrepancies to the true disturbance parameters:

$$\mathcal{L}_{rot} = \frac{1}{N} \sum_{i=1}^{N} \left\| \mathcal{R}_{rot,i} \right\|^2. \tag{21}$$

### C. Simulation Design

To eliminate computational latency anomalies and isolate the convergence behavior of the online-learning PINN estimator, a synchronous fixed-time simulation was utilized for this proof of concept. This guarantees temporal alignment between the plant state and the automatic differentiation tracking within the physics loss. Transitioning these mechanics to an asynchronous, dual-rate Software-in-the-Loop (SIL) framework to

handle real-world computational lag is preserved for future work.

In Algorithm 1, the PINN needs to estimate XY offsets $\hat{d}_x, \hat{d}_y$, and mass $\hat{m}$. To estimate the new system inertia $\hat{\mathbf{J}}$, we computed it from $\hat{d}_x$, $\hat{d}_y$, and $\hat{m}$ combined with the base inertia of the frame, using the parallel axes theorem.

**Algorithm 1** Fixed Time-step Online Parameter Estimation Simulation

**Require:** Simulation duration $T_{max}$, Physics time step $\Delta t$
**Require:** PINN update interval $N$ steps
**Require:** PINN optimization steps $S_{steps}$
**Require:** Replay buffer $\mathcal{B}$ with capacity $W$
**Require:** Initial parameter guesses: $\hat{d}_x, \hat{d}_y, \hat{\mathbf{J}}, \hat{m}$
Init quadrotor physical states $\mathbf{x}_0$
Init PINN $\mathcal{N}_\theta$ with random weights $\theta$
Simulation time $t \leftarrow 0$
Step counter $k \leftarrow 0$
**while** $t \leq T_{max}$ **do**
  Physics engine update state at $t$
  Compute target states $\mathbf{p}_d, \mathbf{v}_d, \mathbf{a}_d, \psi_d, \dot{\psi}_d$ at $t$
  Read current state $\mathbf{x}_k = (\mathbf{p}, \mathbf{v}, \mathbf{R}, \boldsymbol{\omega})$
  Append $\mathbf{x}_k$ and applied rotor thrusts $\mathbf{T}_k$ to buffer $\mathcal{B}$
  **if** $k \pmod N == 0$ **and** $|\mathcal{B}| == W$ **then**
    Compute composite physics loss $\mathcal{L}_{total}$ over $\mathcal{B}$
    Optimize $\theta$ via backpropagation for $S_{steps}$
    Extract updated parameters $\hat{d}_x, \hat{d}_y, \hat{m}$ from $\mathcal{N}_\theta$
    Calculate $\hat{\mathbf{J}}$
  **end if**
  Calculate desired control wrench $\mathbf{F}_{des}, \mathbf{M}_{des}$
  Construct control allocation matrix $\mathbf{A}_{\text{final}}$ using $\hat{d}_x, \hat{d}_y$
  Compute thrust commands $\mathbf{u} \leftarrow \mathbf{A}_{\text{final}}^{\dagger}[\mathbf{F}_{des}^T, \mathbf{M}_{des}^T]^T$
  Apply saturated commands $\mathbf{u}$ to physics engine
  $t \leftarrow t + \Delta t$
  $k \leftarrow k + 1$
**end while**

## V. Validation and Results

### A. Simulation Environment and Setup

Simulations were conducted on a laptop equipped with an NVIDIA RTX 3050 GPU (6 GB VRAM) using Python 3.11.0 with CUDA support. Table II contains the independent variables used in Algorithm 1.

Network optimization is driven by the Adam optimizer with a learning rate of $5 \times 10^{-3}$ [14]. To systematically evaluate the learning characteristics of the PINN, various weighting coefficients ($\lambda$) for the composite physics loss were tested.

The experimental flight sequence initiates at $t = 0$ s with a vertical ascent to a baseline hover. The system stabilizes at an altitude of 3 meters for an initial 5-second window, allowing the physical state to settle before dynamic maneuvers commence. To excite the system's dynamics and test the estimator, the frame is subjected to three distinct spatial trajectories: a figure-8, a circular path, and a uniaxial "lawnmower" motion.

TABLE II
System and Simulation Parameters

| Parameter | Symbol | Value |
|---|---|---|
| Simulation Time-Step | $\Delta t$ | $1/240$ s |
| Gravitational Acceleration | $g$ | 9.81 m/s$^2$ |
| Simulation Duration | $t_{sim}$ | 90 s |
| Base Frame Mass | $m_f$ | 6.8 kg |
| Quadrotor Mass | - | 1.5 kg |
| Payload Mass | $m_p$ | 1.0 kg |
| Payload Volume | $V_p$ | $125 \times 10^{-6}$ m$^3$ |
| Total System Mass | $m_{tot}$ | 7.8 kg |
| Base System Inertia on X axis | $J_{xx}$ | 1.713 |
| Base System Inertia on Y axis | $J_{yy}$ | 1.713 |
| Base System Inertia on Z axis | $J_{zz}$ | 3.413 |
| Max Thrust per Drone | $T_{max}$ | 50 N |
| Unknown Payload Offset | $d_{x,cog}, d_{y,cog}$ | $(-0.1, 0.1)$ m |
| PINN Window Size | $W$ | 50 steps |
| Optimization Frequency | $N_{opt}$ | 5 steps |
| Gradient Steps per Cycle | $S_{grad}$ | 3 steps |

For the sake of brevity in the manuscript, we define only the desired 3D position vector $\mathbf{p}_d(t) = [x_d, y_d, z_d]^T \in \mathbb{R}^3$ for each maneuver. The corresponding feedforward velocity $\mathbf{v}_d(t)$ and acceleration $\mathbf{a}_d(t)$ setpoints utilized by the geometric controller are computed via exact analytical time differentiation of $\mathbf{p}_d(t)$.

The Figure-8 trajectory incorporates a constant radius ($R = 8$ m), an angular frequency ($\omega = 0.10$ rad/s), and a sinusoidal vertical variation ($z_v$) to simultaneously explore translation on frame's XYZ axis:

$$\mathbf{p}_d(t) = \begin{bmatrix} R\sin(\omega t) \\ \frac{R}{2}\sin(2\omega t) \\ z_0 + z_v \sin(\omega t) \end{bmatrix}. \tag{22}$$

The Circular trajectory features a radius ($R = 5$ m) and angular frequency ($\omega = 0.15$ rad/s), retaining the continuous altitude variation:

$$\mathbf{p}_d(t) = \begin{bmatrix} R\sin(\omega t) \\ R(1 - \cos(\omega t)) \\ z_0 + z_v \sin(\omega t) \end{bmatrix}. \tag{23}$$

The Lawnmower maneuver isolates longitudinal dynamics by commanding forward-backward motion exclusively along the world X-axis at an elevated, constant altitude of $z_0 = 5$ m. For a single maneuver phase with an amplitude $A = 2.0$ m and duration $T_{phase} = 5.0$ s, the trajectory is defined as:

$$\mathbf{p}d(t) = \begin{bmatrix} \frac{A}{2}(1 - \cos(\omega_L t)) \\ 0 \\ z_0 \end{bmatrix}, \tag{24}$$

where $\omega_L = \frac{2\pi}{Tphase}$ represents the governing phase angular frequency.

### B. Parameter Convergence and Estimation

The trajectory tracking results illustrated in Fig. 1 demonstrate the significance of integrating physical constraints into the neural network's loss function. As seen in the top row, the data-only baseline struggles to maintain the desired path,

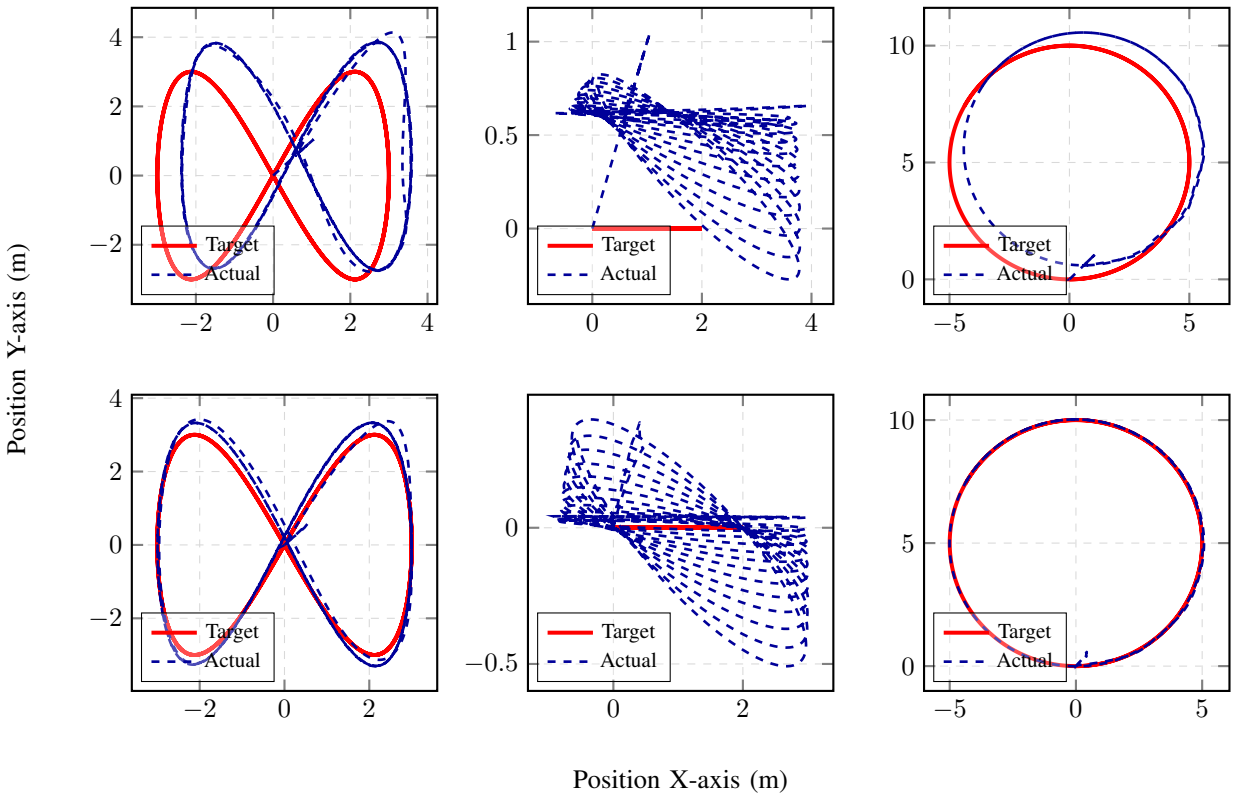



Fig. 1. Trajectory tracking performance in the X-Y plane for each trajectory divided by columns in order Figure-8, lawnmower, and circular path. The top row illustrates the network performance with 0 weight of phyiscs loss, while the bottom row demonstrates the improved trajectory using the proposed method.

TABLE III
STEADY-STATE ESTIMATION ERRORS OVER THE LAST 20 S ACROSS TRAJECTORIES AND $\lambda$ CONFIGURATIONS.

| No | Trajectory | $\lambda_{1,2,3}$ | $\hat{m}$ | $\hat{d}_x$ | $\hat{d}_y$ | $\hat{J}_{xx}$ | $\hat{J}_{yy}$ | $\hat{J}_{zz}$ | Trajectory RMSE (m) |
|---|---|---|---|---|---|---|---|---|---|
| 1 | Figure 8 | 1,1,1 | 7.834 | 0.0489 | 0.0488 | 1.7181 | 1.7180 | 3.4179 | 0.3650 |
| 2 | Figure 8 | 1,0,0 | 6.801 | 0.0000 | 0.0000 | 1.7131 | 1.7131 | 3.4130 | 0.8863 |
| 3 | Figure 8 | 1,1,0 | 7.836 | 0.0000 | 0.0000 | 1.7156 | 1.7156 | 3.4130 | 0.8863 |
| 4 | Figure 8 | 1,0,1 | 6.801 | 1.7180 | 1.7184 | 1.8617 | 1.8610 | 3.7095 | 0.8528 |
| 5 | Figure 8 | 0.1,1,1 | 7.827 | 0.0492 | 0.0490 | 1.7180 | 1.7181 | 3.4180 | 0.3637 |
| 6 | Figure 8 | 10,1,1 | 7.838 | 0.0487 | 0.0487 | 1.7181 | 1.7181 | 3.4180 | 0.3599 |
| 7 | Lawn | 1,1,1 | 7.980 | 0.0433 | 0.0448 | 1.7183 | 1.7189 | 3.4184 | 0.64389 |
| 8 | Lawn | 1,0,0 | 6.801 | 0.0000 | 0.0000 | 1.7131 | 1.7131 | 3.4130 | 1.1394 |
| 9 | Circle | 1,1,1 | 7.770 | 0.0513 | 0.0514 | 1.7180 | 1.7180 | 3.4181 | 0.2315 |
| 10 | Circle | 1,0,0 | 6.801 | 0.0000 | 0.0000 | 1.7131 | 1.7131 | 3.4130 | 0.8528 |

exhibiting significant drift across the Figure-8, lawnmower, and circular maneuvers. Conversely, the bottom row shows that the proposed physics-informed configuration tightly adheres to all three reference trajectories. By embedding the physical system's dynamic model into the loss function, the system successfully compensates the disturbance caused by the payload, resulting in more stable tracking response.

Table III quantifies this performance gap. Relying strictly on the data loss ($\lambda_{1,2,3} = 1, 0, 0$) yields a high trajectory RMSE, reaching 1.1394 m during the lawnmower maneuver and 0.8863 m for the Figure-8 path. Activating the composite physics constraints ($\lambda_{1,2,3} = 1, 1, 1$) substantially reduces these tracking errors to 0.64389 m and 0.3650 m respectively, and down to 0.2315 m for the circular trajectory. Furthermore, Table III reveals the limitation of the data loss only model. It fails to identify the offsets, returning $0.0000$ m for both $\hat{d}_x$ and $\hat{d}_y$ across all trajectories. While the physics-informed approach successfully activates offset estimation to stabilize the system, it does not achieve 100% parameter accuracy. The network finds it challenging to fully close the gap between the estimated and true spatial offsets (converging to roughly $0.05$ m). This indicates that the neural network prioritizes minimizing the overall dynamic residual and tracking error, resulting in a drastically improved flight trajectory without requiring the absolute spatial offset values to be perfectly isolated.

The temporal dynamics of this estimation process are detailed in Fig. 2. The spatial offsets (Offset X and Offset Y)

exhibit convergence during the first 20 seconds of flight across all evaluated trajectories. However, despite this stabilization and the resulting tracking improvements, the offset estimates plateau before reaching the target lines, visually confirming the steady-state gap observed in Table III. In contrast, the system mass ($\hat{m}$) adapts more gradually, steadily converging toward the target value and stabilizing between 40 and 60 seconds. While the lawnmower trajectory induces a slight overshoot in mass estimation, both the Figure-8 and circular paths maintain more stable convergence during the steady-state.

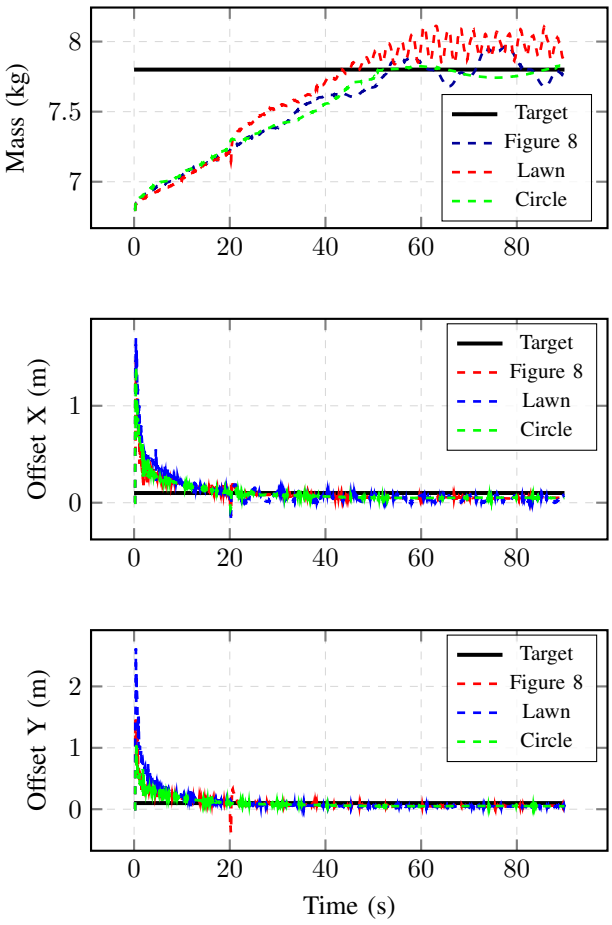



Fig. 2. Parameter convergence for mass, spatial offset X, and spatial offset Y across all evaluated trajectories.

## VI. Conclusion and Future Work

This paper introduced a simulation of online parameter estimation framework for rigid multi-quadrotor payload transportation, driven by a Physics-Informed Neural Network (PINN) estimator. By embedding Euler-Lagrange and Euler-Poincaré equations into the loss function, the network successfully converges on its estimates prooving that the network cannot rely solely on empirical data.

Simulation results validate that integrating this PINN-based estimator into the control allocation matrix enables the system to identify unmodeled payload mass and center-of-mass offsets. The approach effectively arrests transient lateral drift and maintains stable trajectory tracking within a single flight for all tested trajectories. Furthermore, the results indicate that even when absolute parameter convergence is incomplete, the resulting dynamic thrust reallocation is sufficient to restore the baseline controller's authority and counteract unmodeled disturbance torques.

Future work will extend this methodology to software-in-the-loop (SITL) environments by incorporating a publish-subscribe messaging architecture to connect the controller, the simulation, and the online estimator. Ultimately, validating this distributed architecture will serve as the foundational step to transition this framework from simulation to real-world flight testing.

## References


[1] R. Guebsi, S. Mami, and K. Chokmani, "Drones in Precision Agriculture: A Comprehensive Review of Applications, Technologies, and Challenges," *Drones*, vol. 8, no. 11, p. 686, Nov. 2024.

[2] Y. H. Tan, S. Lai, K. Wang, and B. M. Chen, "Cooperative Heavy Lifting Using Unmanned Multi-Agent Systems," in *2018 IEEE 14th International Conference on Control and Automation (ICCA)*. Anchorage, AK: IEEE, Jun. 2018, pp. 1119–1126.

[3] J. Zeng, A. M. Gimenez, E. Vinitsky, J. Alonso-Mora, and S. Sun, "Decentralized Aerial Manipulation of a Cable-Suspended Load using Multi-Agent Reinforcement Learning," Nov. 2025, arXiv:2508.01522 [cs].

[4] B. Peng, B. Su, K. Yi, H. Wu, and D. He, "Modeling and Analysis of Four-UAV-Slung-Load System," in *2023 35th Chinese Control and Decision Conference (CCDC)*. Yichang, China: IEEE, May 2023, pp. 4289–4294.

[5] D. Mellinger, M. Shomin, N. Michael, and V. Kumar, "Cooperative grasping and transport using multiple quadrotors," in *Distributed Autonomous Robotic Systems*. Springer, 2013, pp. 545–558.

[6] H. Ye, H. Yu, B. Liu, J. Han, Y. Fang, and X. Liang, "Nonlinear Robust Control for Dual-UAV Collaborative Aerial Transportation System," in *2023 7th Asian Conference on Artificial Intelligence Technology (ACAIT)*. Jiaxing, China: IEEE, Nov. 2023, pp. 870–875.

[7] C.-A. Lee and T.-H. Cheng, "Concurrent Learning for Cooperative UAV Transportation of Unknown Payloads," *IEEE Open Journal of Control Systems*, vol. 4, pp. 187–198, 2025.

[8] U. Farid, B. Khan, S. M. Ali, and Z. Ullah, "A Digital Twin Model for UAV Control to Lift Irregular-Shaped Payloads Using Robust Model Predictive Control," *Machines*, vol. 13, no. 11, p. 1069, Nov. 2025.

[9] M. Raissi, P. Perdikaris, and G. E. Karniadakis, "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations," *Journal of Computational physics*, vol. 378, pp. 686–707, 2019.

[10] R. Mojgani, M. Balajewicz, and P. Hassanzadeh, "Kolmogorov n–width and Lagrangian physics-informed neural networks: A causality-conforming manifold for convection-dominated PDEs," *Computer Methods in Applied Mechanics and Engineering*, vol. 404, p. 115810, Feb. 2023.

[11] F. Furrer, M. Burri, M. Achtelik, and R. Siegwart, *Robot Operating System (ROS): The Complete Reference (Volume 1)*. Cham: Springer International Publishing, 2016, ch. RotorS—A Modular Gazebo MAV Simulator Framework, pp. 595–625.

[12] S. Lakshminarayana, S. Sthapit, and C. Maple, "Application of Physics-Informed Machine Learning Techniques for Power Grid Parameter Estimation," *Sustainability*, vol. 14, no. 4, p. 2051, Feb. 2022.

[13] M. Raissi, P. Perdikaris, and G. Karniadakis, "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations," *Journal of Computational Physics*, vol. 378, pp. 686–707, Feb. 2019.

[14] G. Serrano, M. Jacinto, J. Ribeiro-Gomes, J. Pinto, B. J. Guerreiro, A. Bernardino, and R. Cunha, "Physics-Informed Neural Network for Multirotor Slung Load Systems Modeling," May 2024, arXiv:2405.09428 [cs].

[15] G. Haghighatdoost, "Euler Poincare Dynamics and Control on Lie Groupoids," May 2026, arXiv:2509.20044 [math.DS].